

# Advanced Algorithms: The 20/80 Summary

COMS3005A - University of the Witwatersrand

2024

## Introduction

This comprehensive summary distills the core principles (20%) that enable deep understanding of all algorithmic concepts (80%) from the Advanced Analysis of Algorithms course. The 20% represents the fundamental mental models and ways of thinking that allow you to analyze, design, and reason about any algorithm, while the 80% represents the specific implementations and applications of these principles. Master these fundamentals to develop true algorithmic intuition.

## 1 The 20%: Foundational Principles

### 1.1 The Trade-Off Triad: The Fundamental Constraint

- **Time vs Space vs Optimality:** Every algorithmic decision represents a point in this three-dimensional trade-off space. Understanding where your solution lies in this space is crucial for making informed design decisions.
- **Time Complexity:** How the runtime scales with input size. This is typically the primary concern, but not always. Some systems prioritize predictable performance over absolute speed.
- **Space Complexity:** How memory usage scales with input size. In memory-constrained environments (embedded systems, mobile devices), this can be the dominant constraint.
- **Optimality:** Whether the solution is guaranteed to be the best possible. In practice, we often accept suboptimal solutions that are "good enough" but much faster to compute.
- **Real-world Example:** Merge Sort provides optimal  $O(n \log n)$  time but requires  $O(n)$  extra space. Quick Sort provides  $O(n \log n)$  average time with only  $O(\log n)$  extra space but has  $O(n^2)$  worst case. The choice depends on which trade-offs are acceptable for your specific context.

### 1.2 Asymptotic Analysis: The Language of Scalability

- **Big-O ( $O$ ):** Describes the *upper bound* of growth. When we say an algorithm is  $O(n^2)$ , we mean it will *not grow faster than* quadratic time. This is about *worst-case* guarantees.
- **Omega ( $\Omega$ ):** Describes the *lower bound* of growth.  $\Omega(n \log n)$  means the algorithm requires *at least* linearithmic time. This captures *best-case* requirements.
- **Theta ( $\Theta$ ):** Describes a *tight bound* where the algorithm grows *exactly* at this rate (both upper and lower bounds match). This is the most informative notation when applicable.
- **Practical Interpretation:**
  - $O(1)$ : Instant, regardless of input size (hash table lookup)
  - $O(\log n)$ : Incredibly fast even for huge inputs (binary search)

- $O(n)$ : Proportional time (linear scan)
- $O(n \log n)$ : The "bar" for efficient algorithms (optimal sorting)
- $O(n^2)$ : Becomes slow for moderate inputs (naïve sorting)
- $O(2^n)$ : Impractical except for tiny inputs (brute force)
- **Why Constants Don't Matter:** For large  $n$ , the growth rate dominates any constant factors.  $1000n$  is eventually better than  $2n^2$  because linear growth will always beat quadratic growth for sufficiently large  $n$ .

### 1.3 Problem-Solving Paradigms: The Algorithm Designer's Toolkit

- **Brute Force:** The simplest approach - enumerate all possibilities. While often impractical ( $O(n!)$ ,  $O(2^n)$ ), it serves as:
  - A baseline for comparing more sophisticated algorithms
  - A solution for very small input sizes
  - A starting point for understanding the problem structure
- **Decrease & Conquer:** Solve a smaller instance of the *same* problem, then extend the solution. Examples:
  - **Insertion Sort:** Sort first  $n - 1$  elements, then insert the  $n^{th}$  element
  - **Binary Search:** Eliminate half the search space at each step
- **Divide & Conquer:** The most powerful paradigm:
  - **Divide:** Split problem into independent subproblems
  - **Conquer:** Solve subproblems recursively
  - **Combine:** Merge solutions to form final answer
  - **Examples:** Merge Sort, Quick Sort, Closest Pair
- **Transform & Conquer:** Convert the problem into a different representation where it becomes easier:
  - Sort data first to enable efficient searching
  - Use hash functions to transform complex keys into simple indices
  - Represent graphs as adjacency matrices vs lists based on operations needed
- **Greedy:** Make the locally optimal choice at each step. The key insight is recognizing when local optimality guarantees global optimality.
- **Dynamic Programming:** The art of solving each subproblem exactly once by storing solutions. The essence is overlapping subproblems + optimal substructure.

### 1.4 Data Structures as Algorithm Enablers

- **Critical Insight:** Algorithms don't exist in isolation - they're built on data structures that enable specific operations efficiently.
- **Array:**
  - **Strengths:**  $O(1)$  random access, cache-friendly, simple

- **Weaknesses:** Fixed size, slow insertion/deletion ( $O(n)$ )
- **Best for:** Sequential processing, known size datasets
- **Linked List:**
  - **Strengths:**  $O(1)$  insertion/deletion at known positions, dynamic size
  - **Weaknesses:**  $O(n)$  random access, poor cache performance
  - **Best for:** Frequent insertions/deletions, unknown size data
- **Hash Table:**
  - **Strengths:**  $O(1)$  average case operations, flexible keys
  - **Weaknesses:**  $O(n)$  worst case, no ordering, hash function dependent
  - **Best for:** Dictionary operations, membership testing
- **Binary Search Trees:**
  - **Strengths:**  $O(\log n)$  operations, maintains ordering, range queries
  - **Weaknesses:**  $O(n)$  worst case if unbalanced, more complex
  - **Best for:** Ordered data, range queries, sequential access

## 1.5 Correctness & Optimality: Mathematical Rigor

- **Proof Techniques:**
  - **Induction:** Perfect for recursive algorithms and loop invariants
  - **Contradiction:** Assume algorithm is wrong and derive contradiction
  - **Adversary Arguments:** Construct worst-case inputs that force minimum work
  - **Loop Invariants:** Properties that hold before/during/after loop execution
- **Lower Bounds:** The theoretical minimum work required to solve a problem:
  - **Comparison Sorting:**  $\Omega(n \log n)$  from decision tree height
  - **Searching Unsorted Data:**  $\Omega(n)$  from adversary argument
  - **Importance:** Tells us when we can stop looking for better algorithms
- **Optimality:** An algorithm is optimal if its complexity matches the problem's lower bound. Merge Sort is optimal for comparison sorting. Binary Search is optimal for searching sorted arrays.

# 2 Searching Algorithms

## 2.1 Linear Search: The Baseline

- **Fundamental Operation:** Sequential comparison of each element with the key until found or end reached.
- **Complexity Analysis:**
  - **Best Case:**  $\Theta(1)$  - key is first element
  - **Worst Case:**  $\Theta(n)$  - key is last element or not present

- **Average Case:**  $\Theta(n)$  - key is equally likely in any position
- **Optimality Proof:** For unordered data, any correct algorithm must examine every element in the worst case. An adversary could always place the key in the last position examined.
- **Variants:**
  - **Sentinel Technique:** Place key at end to eliminate bound checking
  - **Probability-based:** Search more likely positions first
- **Real-world Use:** Small datasets, unsorted data, when simplicity is more important than speed.

## 2.2 Binary Search: Divide and Conquer in Action

- **Prerequisite:** Data must be sorted - this is the "transform" step that enables the efficient "conquer" step.
- **Algorithm Intuition:** At each step, the search space is halved. This exponential reduction is what gives logarithmic complexity.
- **Complexity Analysis:**
  - **Recurrence Relation:**  $T(n) = T(n/2) + O(1)$
  - **Solving:**  $T(n) = T(n/4) + 2 = T(n/8) + 3 = \dots = T(1) + \log n = O(\log n)$
  - This represents the height of the decision tree
- **Implementation Details:**
  - **Midpoint Calculation:** Use  $\lfloor low + (high - low)/2 \rfloor$  to avoid overflow
  - **Termination Condition:**  $low \leq high$  vs  $low < high$  affects edge cases
  - **Element Not Found:** Return insertion point for complete specification
- **Variants:**
  - **Lower/Upper Bound:** Find first/last occurrence in duplicates
  - **Exponential Search:** Unknown array size
  - **Interpolation Search:**  $O(\log \log n)$  for uniform distributions
- **Optimality:** The decision tree argument shows that any comparison-based search on sorted data requires  $\Omega(\log n)$  comparisons.

## 3 Sorting Algorithms

### 3.1 Comparison Sort Lower Bound: A Fundamental Limit

- **The Core Argument:**
  - There are  $n!$  possible permutations of  $n$  elements
  - Each comparison gives at most 1 bit of information (true/false)
  - To distinguish among  $n!$  permutations, we need  $\log(n!)$  bits
  - $\log(n!) \approx n \log n - n \log e + \Theta(\log n)$  by Stirling's approximation
  - Therefore,  $\Omega(n \log n)$  comparisons are required

- **Decision Tree Model:** Any comparison-based sort corresponds to a decision tree where:
  - Internal nodes represent comparisons
  - Leaves represent sorted permutations
  - The height of the tree is the number of comparisons needed
- **Implications:**
  - Merge Sort, Heap Sort are optimal (they achieve  $O(n \log n)$ )
  - Quick Sort is optimal on average ( $O(n \log n)$  expected)
  - No comparison-based sort can do better than  $O(n \log n)$
- **Non-comparison Sorts:** Bucket Sort, Radix Sort can achieve  $O(n)$  by exploiting additional information about the data.

### 3.2 Quadratic Sorts: When Simple is Better

- **Selection Sort:**
  - **Mechanism:** Repeatedly find minimum element and swap to front
  - **Complexity:** Always  $\Theta(n^2)$  comparisons,  $O(n)$  swaps
  - **Advantage:** Minimal data movement - good when writes are expensive
  - **Stability:** Not stable due to swapping distant elements
- **Bubble Sort:**
  - **Mechanism:** Repeatedly swap adjacent inverted elements
  - **Complexity:**  $O(n^2)$  worst/average,  $O(n)$  best (already sorted)
  - **Optimization:** Stop if no swaps in a pass (detects sorted array)
  - **Advantage:** Simple to implement, stable, detects sorted input
  - **Real Use:** Educational purposes, tiny datasets
- **Insertion Sort:**
  - **Mechanism:** Build sorted array one element at a time by insertion
  - **Complexity:**  $O(n^2)$  worst/average,  $O(n)$  best (already sorted)
  - **Advantages:**
    - \* Excellent for small  $n$  ( $n \leq 50$ )
    - \* Adaptive: very fast on nearly sorted data
    - \* Stable: preserves order of equal elements
    - \* In-place:  $O(1)$  extra space
  - **Real Use:** Small arrays, as the base case for hybrid sorts like Timsort
- **When to Use Quadratic Sorts:**
  - Small datasets where  $n^2$  is acceptable
  - Nearly sorted data (Insertion Sort shines here)

- Simple implementation is more important than absolute speed
- Educational contexts to understand sorting fundamentals

### 3.3 Merge Sort: The Optimal Workhorse

- **Divide and Conquer Structure:**

- **Divide:** Split array into two equal halves
- **Conquer:** Recursively sort both halves
- **Combine:** Merge the two sorted halves

- **Merge Operation:** The key to efficiency:

- Compare elements from front of both subarrays
- Copy smaller element to result
- Continue until one subarray is exhausted
- Copy remaining elements
- Time:  $\Theta(n)$  for merging two subarrays of total size  $n$

- **Complexity Analysis:**

- **Recurrence:**  $T(n) = 2T(n/2) + \Theta(n)$
- **Solving:**

$$\begin{aligned}
 T(n) &= 2T(n/2) + cn \\
 &= 2[2T(n/4) + c(n/2)] + cn = 4T(n/4) + 2cn \\
 &= 4[2T(n/8) + c(n/4)] + 2cn = 8T(n/8) + 3cn \\
 &\vdots \\
 &= 2^k T(n/2^k) + kc n
 \end{aligned}$$

- When  $n/2^k = 1$ ,  $k = \log n$ , so  $T(n) = nT(1) + cn \log n = \Theta(n \log n)$

- **Properties:**

- **Stable:** Yes (if merge prefers left element on ties)
- **Parallelizable:** Easy to parallelize the recursive calls
- **External Sorting:** Can sort data too large for memory
- **Disadvantage:**  $O(n)$  extra space requirement

- **Variants:**

- **Natural Merge Sort:** Exploits existing sorted runs
- **Bottom-up Merge Sort:** Iterative version avoids recursion
- **In-place Merge Sort:** Complex but reduces space to  $O(1)$

### 3.4 Quick Sort: The Practical Champion

- **Algorithm Structure:**

- **Choose Pivot:** Select an element to partition around

- **Partition:** Rearrange so elements  $<$  pivot come before, elements  $>$  pivot come after
- **Recurse:** Sort left and right partitions
- **Partition Strategies:**
  - **Lomuto:** Simpler but less efficient, multiple swaps
  - **Hoare:** More efficient, fewer swaps, complex invariants
  - **Dutch National Flag:** Handles duplicates efficiently
- **Pivot Selection Critical:**
  - **Worst Case:** Already sorted data with first/last pivot ( $O(n^2)$ )
  - **Best Case:** Median element as pivot ( $O(n \log n)$ )
  - **Good Strategies:**
    - \* **Median-of-Three:** First, middle, last - choose median
    - \* **Random:** Provides probabilistic guarantees
    - \* **Introselect:** Hybrid for guaranteed  $O(n \log n)$
- **Complexity Analysis:**
  - **Worst Case:**  $O(n^2)$  when pivot is always min/max
  - **Best Case:**  $O(n \log n)$  when pivot is always median
  - **Average Case:**  $O(n \log n)$  with constant factor  $\approx 1.39n \log n$
  - **Expected Case:**  $O(n \log n)$  with randomized pivot
- **Why Quick Sort Wins in Practice:**
  - Excellent cache performance (sequential access pattern)
  - Small constant factors
  - In-place ( $O(\log n)$ ) stack space for recursion
  - Hardware-friendly access patterns
- **Optimizations:**
  - **Hybrid Approach:** Switch to Insertion Sort for small subarrays
  - **Tail Recursion:** Eliminate one recursive call
  - **Three-way Partitioning:** Handle duplicates efficiently

## 4 Informed Search & Heuristics

### 4.1 A\* Search: The Perfect Balance

- **The Insight:** Best-first search that considers both:
  - $g(n)$ : Actual cost from start to node  $n$  (like Uniform Cost Search)
  - $h(n)$ : Estimated cost from node  $n$  to goal (like Greedy Search)
- **Evaluation Function:**  $f(n) = g(n) + h(n)$  represents estimated total cost through node  $n$

- **Optimality Conditions:**

- **Admissible Heuristic:**  $h(n) \leq h^*(n)$  for all  $n$  (never overestimates)
- **Consistent Heuristic:**  $h(n) \leq c(n, n') + h(n')$  for all  $n, n'$  (triangle inequality)
- **Tree Search:** Optimal with admissible heuristic
- **Graph Search:** Optimal with consistent heuristic

- **Why Consistency Matters:** Ensures  $f(n)$  is non-decreasing along any path, so we never need to reconsider nodes

- **Complexity:**

- **Time:**  $O(b^{\epsilon d})$  where  $\epsilon$  is heuristic error
- **Space:**  $O(b^d)$  - stores all generated nodes
- **Optimal Efficiency:** No other optimal algorithm expands fewer nodes

- **Implementation:**

- **Priority Queue:** Ordered by  $f(n) = g(n) + h(n)$
- **Closed Set:** Track visited nodes for graph search
- **Path Reconstruction:** Store parent pointers

## 4.2 Heuristic Design: The Art of Estimation

- **Admissible Heuristics for Pathfinding:**

- **Manhattan Distance:**  $|x_1 - x_2| + |y_1 - y_2|$  for grid movement
- **Euclidean Distance:**  $\sqrt{(x_1 - x_2)^2 + (y_1 - y_2)^2}$  for direct movement
- **Chebyshev Distance:**  $\max(|x_1 - x_2|, |y_1 - y_2|)$  for 8-direction movement

- **Relaxation Method:** Create admissible heuristics by solving a relaxed version of the problem:

- **8-puzzle:** Remove wall constraints = Manhattan distance
- **8-puzzle:** Remove tile interaction = misplaced tiles count
- The more constraints removed, the easier to compute but less informative

- **Pattern Databases:** Precompute solutions to subproblems:

- Store optimal costs for pattern configurations
- Use as lookup during search
- Very effective for puzzles like 15-puzzle, Rubik's cube

- **Learning Heuristics:**

- Machine learning to predict costs
- Neural networks for complex state evaluation
- Requires training data but can discover non-obvious patterns

- **Heuristic Quality:** Measured by:

- **Effective Branching Factor:** How much the heuristic reduces search
- **Informedness:** How close  $h(n)$  is to  $h^*(n)$  without overestimating
- The perfect heuristic  $h(n) = h^*(n)$  would solve the problem instantly

## 5 Greedy Algorithms

### 5.1 The Greedy Choice Property

- **Core Idea:** A series of locally optimal choices leads to a globally optimal solution.
- **When It Works:**
  - **Matroid Structure:** Problems that can be represented as matroids
  - **Exchange Argument:** Can transform any solution to the greedy solution without making it worse
  - **Greedy Stays Ahead:** The greedy solution is never worse than any partial solution
- **Optimal Substructure:** An optimal solution contains optimal solutions to subproblems. This is shared with Dynamic Programming, but greedy makes irrevocable choices.
- **Proof Techniques:**
  - **Greedy Stays Ahead:** Show greedy is always at least as good as any other solution at each step
  - **Exchange Argument:** Transform any optimal solution into the greedy solution
  - **Matroid Theory:** Formal mathematical framework for greedy algorithms

### 5.2 Classic Greedy Problems

- **Optimal Service (Scheduling):**
  - **Problem:** Minimize total waiting time for  $n$  customers
  - **Greedy Strategy:** Serve shortest job first
  - **Proof:** Exchange argument - swapping two jobs shows optimality
  - **Complexity:**  $O(n \log n)$  for sorting
- **Change Making:**
  - **Problem:** Make change with fewest coins
  - **Greedy Strategy:** Use largest denomination possible
  - **When Optimal:** For "canonical" coin systems (1, 5, 10, 25, ...)
  - **Counterexample:** Coins 1, 3, 4, target 6: greedy gives  $4+1+1=3$  coins, optimal is  $3+3=2$  coins
  - **Characterization:** Coin system is canonical if greedy works for all amounts
- **Huffman Coding:**
  - **Problem:** Optimal prefix-free compression
  - **Greedy Strategy:** Merge least frequent symbols
  - **Optimality:** Exchange argument shows no better code exists

- **Complexity:**  $O(n \log n)$  with priority queue
- **Interval Scheduling:**
  - **Problem:** Schedule maximum number of non-overlapping intervals
  - **Greedy Strategy:** Choose interval with earliest finish time
  - **Proof:** Greedy stays ahead - always leaves maximum remaining capacity
- **Minimum Spanning Tree:**
  - **Problem:** Connect all vertices with minimum total edge weight
  - **Greedy Strategies:**
    - \* **Kruskal's:** Add smallest edge that doesn't create cycle
    - \* **Prim's:** Grow tree from vertex adding cheapest connecting edge
  - **Optimality:** Cut property guarantees optimality

### 5.3 When Greedy Fails

- **Typical Failure Modes:**
  - Local optimum doesn't lead to global optimum
  - Decisions are not reversible
  - Problem lacks optimal substructure
- **Examples of Failure:**
  - **0-1 Knapsack:** Greedy by value/weight ratio fails
  - **Traveling Salesman:** Nearest neighbor heuristic can be arbitrarily bad
  - **Non-canonical Coin Systems:** Greedy change making suboptimal
- **Testing Greedy Approaches:**
  - Try to construct counterexamples
  - Check if exchange argument works
  - Verify greedy choice property holds

## 6 Dynamic Programming

### 6.1 Recognizing DP Problems

- **Overlapping Subproblems:** The same subproblem is solved multiple times in a naive recursive solution.
- **Optimal Substructure:** An optimal solution can be constructed from optimal solutions of subproblems.
- **Common Patterns:**
  - "Find the minimum/maximum cost/path"
  - "Count the number of ways to..."
  - "Decide if possible to achieve..."

- Sequences, strings, grids, trees, graphs
- **DP vs Divide & Conquer:**
  - **Divide & Conquer:** Subproblems are independent
  - **Dynamic Programming:** Subproblems overlap
- **DP vs Greedy:**
  - **Greedy:** Make choice and solve one subproblem
  - **Dynamic Programming:** Try all choices and combine results

## 6.2 DP Implementation Strategies

- **Top-Down with Memoization:**
  - Write natural recursive solution
  - Add cache to store computed results
  - Check cache before computing
  - Store result in cache after computing
  - **Advantages:** Natural, computes only needed subproblems
  - **Disadvantages:** Recursion overhead, harder to optimize space
- **Bottom-Up with Tabulation:**
  - Identify dependency order of subproblems
  - Solve smallest subproblems first
  - Build up to original problem
  - **Advantages:** No recursion overhead, easier space optimization
  - **Disadvantages:** May compute unnecessary subproblems, less intuitive
- **State Definition:** The art of DP is defining the state:
  - What parameters define a subproblem?
  - What is the recurrence relation between states?
  - What are the base cases?
  - How do we reconstruct the solution?

## 6.3 Classic DP Problems and Patterns

- **Fibonacci Sequence:**
  - **Naive:**  $O(2^n)$  - exponential blowup
  - **DP:**  $O(n)$  time,  $O(n)$  space
  - **Optimized:**  $O(n)$  time,  $O(1)$  space - only need last two values
  - **Pattern:** Caching overlapping computations
- **Change Making:**

- **Problem:** Minimum coins to make amount with given denominations
  - **State:**  $dp[i] = \min$  coins for amount  $i$
  - **Recurrence:**  $dp[i] = \min(dp[i - c_j] + 1)$  for all coins  $c_j \leq i$
  - **Complexity:**  $O(n \cdot k)$  for amount  $n$ ,  $k$  coin types
  - **Advantage:** Handles non-canonical coin systems
- **Longest Common Subsequence:**
    - **State:**  $dp[i][j] = \text{LCS of first } i \text{ chars of A and first } j \text{ chars of B}$
    - **Recurrence:**
      - \* If  $A[i] = B[j]$ :  $dp[i][j] = dp[i - 1][j - 1] + 1$
      - \* Else:  $dp[i][j] = \max(dp[i - 1][j], dp[i][j - 1])$
    - **Complexity:**  $O(mn)$  for strings of length  $m, n$
    - **Reconstruction:** Trace back through DP table
  - **Matrix Chain Multiplication:**
    - **Problem:** Optimal parenthesization of matrix multiplication
    - **State:**  $dp[i][j] = \min \text{ cost to multiply matrices } i \text{ through } j$
    - **Recurrence:**  $dp[i][j] = \min(dp[i][k] + dp[k + 1][j] + \text{cost})$  for  $i \leq k < j$
    - **Complexity:**  $O(n^3)$  vs brute force  $O(2^n)$
  - **Knapsack Problems:**
    - **0-1 Knapsack:** Each item take or leave
    - **Unbounded Knapsack:** Unlimited copies of each item
    - **State:**  $dp[i][w] = \max \text{ value with first } i \text{ items and weight } w$
    - **Complexity:**  $O(nW)$  pseudo-polynomial
  - **Edit Distance:**
    - **Problem:** Minimum operations to transform string A to B
    - **Operations:** Insert, delete, replace
    - **State:**  $dp[i][j] = \text{edit distance of first } i \text{ chars of A and first } j \text{ chars of B}$
    - **Recurrence:** Consider all three operations

## 6.4 DP Optimization Techniques

- **Space Optimization:**
  - Reuse arrays when only previous row/column needed
  - Fibonacci:  $O(1)$  space instead of  $O(n)$
  - Knapsack:  $O(W)$  space instead of  $O(nW)$
- **State Reduction:**
  - Find more compact state representation

- Use bitmasking for small sets
- Exploit symmetries in the problem
- **Decision Optimization:**
  - Monotonic queue/stack for certain recurrences
  - Convex hull trick for specific cost functions
  - Divide and conquer optimization

## 7 Computational Geometry

### 7.1 Closest Pair Problem: Divide and Conquer Mastery

- **Problem Statement:** Given  $n$  points in the plane, find the pair with smallest Euclidean distance.
- **Brute Force:** Check all  $\binom{n}{2}$  pairs -  $O(n^2)$
- **Divide and Conquer Approach:**
  1. **Sort by x-coordinate:**  $O(n \log n)$  preprocessing
  2. **Divide:** Split points into left and right halves by x-coordinate
  3. **Conquer:** Recursively find closest pairs in left and right
  4. **Combine:** Check pairs that cross the dividing line
- **The Key Insight - The Strip:**
  - Let  $\delta = \min(\delta_{left}, \delta_{right})$
  - Only need to consider points within  $\delta$  of the dividing line
  - For each point in the strip, only check next 7 points by y-coordinate
- **Why 7 Points? Geometric Proof:**
  - Consider a  $\delta \times 2\delta$  rectangle centered on the dividing line
  - Divide into  $8 \delta/2 \times \delta/2$  squares
  - Each square can contain at most 1 point (otherwise  $\delta$  would be smaller)
  - Therefore, at most 8 points in the rectangle
  - For any point, only need to check the other 7
- **Complexity Analysis:**
  - **Recurrence:**  $T(n) = 2T(n/2) + O(n)$
  - **Solving:** By master theorem,  $T(n) = O(n \log n)$
  - **Total:**  $O(n \log n)$  for sort +  $O(n \log n)$  for algorithm =  $O(n \log n)$
- **Implementation Details:**
  - **Sorting:** Pre-sort by x and maintain y-order during recursion
  - **Merge by y:**  $O(n)$  merge of two sorted y-lists
  - **Strip Processing:**  $O(n)$  for processing all points in strip

- **Generalization:**

- **Higher Dimensions:** Becomes  $O(n \log^{d-1} n)$  for  $d$  dimensions
- **Other Metrics:** Works for any  $L_p$  metric
- **Approximate Versions:** Can find  $(1 + \epsilon)$ -approximation faster

## 7.2 Other Geometric Algorithms

- **Convex Hull:**

- **Graham Scan:**  $O(n \log n)$  - sort by angle and scan
- **Jarvis March:**  $O(nh)$  - output-sensitive, good for small hulls
- **QuickHull:**  $O(n \log n)$  expected - divide and conquer

- **Line Segment Intersection:**

- **Sweep Line Algorithm:**  $O(n \log n)$
- **Sweep vertical line**, maintain active segments in balanced BST
- Detect intersections when sweep line reaches endpoints

- **Point Location:**

- **Preprocessing:** Build data structure for point-in-polygon queries
- **Trapezoidal Decomposition:**  $O(n \log n)$  preprocessing,  $O(\log n)$  query

- **Voronoi Diagrams & Delaunay Triangulation:**

- **Fortune's Algorithm:**  $O(n \log n)$  sweep line
- **Applications:** Nearest neighbor, mesh generation, terrain modeling

## 8 Data Structures

### 8.1 Hash Tables: The Power of Direct Access

- **Core Idea:** Use a hash function to map keys to array indices, enabling  $O(1)$  average-case operations.

- **Hash Function Design:**

- **Requirements:** Fast to compute, uniform distribution, deterministic
- **Good Hash Functions:** MurmurHash, CityHash, cryptographic hashes
- **Universal Hashing:** Family of hash functions for worst-case guarantees

- **Collision Resolution Strategies:**

- **Separate Chaining:**
  - \* Each bucket is a linked list
  - \* Simple, handles arbitrary load factors
  - \* Cache-unfriendly, pointer overhead
- **Open Addressing:**

- \* Store all entries in the table itself
- \* **Linear Probing:**  $h(k, i) = (h(k) + i) \bmod m$
- \* **Quadratic Probing:**  $h(k, i) = (h(k) + c_1i + c_2i^2) \bmod m$
- \* **Double Hashing:**  $h(k, i) = (h_1(k) + i \cdot h_2(k)) \bmod m$

- **Load Factor and Performance:**

- $\alpha = n/m$  where  $n$  = entries,  $m$  = buckets
- **Separate Chaining:** Expected chain length =  $\alpha$
- **Open Addressing:** Performance degrades as  $\alpha \rightarrow 1$
- **Rehashing:** Double table size when  $\alpha$  exceeds threshold

- **Complexity Analysis:**

- **Average Case:**  $O(1)$  for all operations with good hash function
- **Worst Case:**  $O(n)$  if all keys hash to same bucket
- **Amortized:**  $O(1)$  considering periodic rehashing

- **Advanced Variants:**

- **Cuckoo Hashing:** Multiple hash functions, guaranteed  $O(1)$  worst-case lookup
- **Perfect Hashing:**  $O(1)$  worst-case for static sets
- **Bloom Filters:** Space-efficient probabilistic membership testing

## 8.2 Balanced BSTs: Red-Black Trees

- **Motivation:** Maintain  $O(\log n)$  operations while preserving order, unlike hash tables.

- **Red-Black Properties:**

1. Every node is either red or black
2. The root is always black
3. All leaves (NIL) are black
4. If a node is red, then both its children are black (no two consecutive reds)
5. Every path from a node to any of its descendant NIL nodes contains the same number of black nodes

- **Consequences of Properties:**

- **Height Bound:**  $h \leq 2 \log(n + 1)$  - approximately balanced
- **Black Height:** Same for all paths, ensures balance
- **Longest Path:** At most twice the shortest path

- **Rotation Operations:**

- **Left Rotation:** Make right child the new root
- **Right Rotation:** Make left child the new root
- Preserve BST property while changing tree structure

- **Insertion Cases:**

- **Case 1:** Uncle is red - recolor parent, uncle, grandparent
- **Case 2:** Uncle is black, node is right child - rotate to make left child
- **Case 3:** Uncle is black, node is left child - rotate grandparent
- At most 2 rotations needed for insertion

- **Deletion Cases:**

- More complex due to "double black" nodes
- 4 main cases with symmetric variants
- May require rotations and recoloring up to the root

- **Comparison with Other Balanced BSTs:**

- **AVL Trees:** Stricter balance, faster lookups, slower insert/delete
- **B-Trees:** Better for disk-based systems, higher branching factor
- **Splay Trees:** Amortized bounds, no balance guarantees, good locality

- **Real-world Use:**

- **Java:** TreeMap, TreeSet
- **C++:** std::map, std::set (typically red-black)
- **Linux Kernel:** Completely fair scheduler uses red-black trees

## 9 Adversarial Search

### 9.1 Minimax Algorithm: Optimal Play

- **Problem Setting:** Two-player zero-sum games with perfect information.

- **Key Assumptions:**

- Both players play optimally
- Players have opposite objectives (zero-sum)
- All information is visible to both players

- **Minimax Values:**

- $V(s)$ : Utility of state  $s$  for MAX player
- If  $s$  is terminal:  $V(s) = \text{known utility}$
- If MAX to move:  $V(s) = \max_{s' \in \text{successors}(s)} V(s')$
- If MIN to move:  $V(s) = \min_{s' \in \text{successors}(s)} V(s')$

- **Algorithm:**

- Recursively compute minimax values from leaves up
- MAX chooses move that maximizes minimax value
- MIN chooses move that minimizes minimax value

- **Complexity:**

- **Time:**  $O(b^m)$  where  $b$  = branching factor,  $m$  = maximum depth
- **Space:**  $O(bm)$  for DFS implementation
- **Example:** Chess has  $b \approx 35$ ,  $m \approx 100$ , making  $35^{100}$  infeasible

- **Evaluation Functions:**

- Estimate utility of non-terminal states
- Should correlate with actual probability of winning
- Often weighted linear functions of features
- Example: Chess - material advantage, piece activity, king safety

- **Depth-Limited Search:**

- Search to fixed depth rather than leaves
- Use evaluation function at depth limit
- Must handle quiescence - avoid evaluating volatile positions

## 9.2 Alpha-Beta Pruning: Optimizing Minimax

- **Core Idea:** Prune branches that cannot affect the final decision.

- **Alpha ( $\alpha$ ):** Best value MAX can guarantee so far (lower bound)

- **Beta ( $\beta$ ):** Best value MIN can guarantee so far (upper bound)

- **Pruning Conditions:**

- **MAX node:** Prune if value  $\geq \beta$  (MIN won't allow this path)
- **MIN node:** Prune if value  $\leq \alpha$  (MAX has better option)

- **Algorithm:**

- Initialize:  $\alpha = -\infty$ ,  $\beta = +\infty$
- MAX nodes: Update  $\alpha$ , prune if value  $\geq \beta$
- MIN nodes: Update  $\beta$ , prune if value  $\leq \alpha$

- **Effectiveness:**

- **Best Case:**  $O(b^{m/2})$  - effectively doubles search depth
- **Worst Case:**  $O(b^m)$  - no pruning occurs
- **Average Case:**  $O(b^{3m/4})$  for reasonable move ordering

- **Move Ordering:**

- Critical for effective pruning
- Try best moves first (captures, threats, good positional moves)
- Use iterative deepening to inform move ordering
- Killer heuristic: moves that were good in similar positions

- **Enhancements:**

- **Iterative Deepening:** Search to depth 1, 2, 3, ... using previous results
- **Transposition Tables:** Cache previously computed positions
- **Null Move Pruning:** Assume passing is worse than any real move
- **Quiescence Search:** Extend search until position is stable
- **Real-world Performance:**
  - Modern chess engines search 15-20 plies deep
  - Alpha-beta crucial for making deep search feasible
  - Combined with sophisticated evaluation functions and opening books

## 10 P vs NP Complexity

### 10.1 Complexity Classes Overview

- **P (Polynomial Time):** Decision problems solvable in polynomial time by deterministic Turing machines.
  - Examples: Sorting, shortest path, minimum spanning tree
  - Considered "efficiently solvable" in practice
- **NP (Nondeterministic Polynomial Time):** Decision problems where "yes" answers can be verified in polynomial time.
  - Examples: SAT, traveling salesman, graph coloring
  - Can be solved in polynomial time by nondeterministic Turing machines
- **NP-complete:** The hardest problems in NP. If any NP-complete problem is in P, then P = NP.
  - Examples: SAT, 3-SAT, vertex cover, Hamiltonian path
  - All NP-complete problems are polynomially reducible to each other
- **NP-hard:** Problems at least as hard as NP-complete problems, but not necessarily in NP.
  - Examples: Halting problem, chess optimal play
  - May be harder than NP-complete
- **The P vs NP Question:** Is every problem that can be verified quickly also solvable quickly?
  - One of the seven Millennium Prize Problems
  - Most experts believe  $P \neq NP$
  - Would have profound implications for cryptography, optimization, AI

### 10.2 Practical Implications

- **Algorithm Design Strategy:**
  - First try to find polynomial-time algorithm
  - If problem appears hard, check if it's NP-complete
  - For NP-complete problems, consider:

- \* Approximation algorithms
- \* Heuristics and local search
- \* Parameterized algorithms
- \* Exact algorithms for small instances

- **Approximation Algorithms:**

- Polynomial-time algorithms with guaranteed solution quality
- Example: 2-approximation for vertex cover
- Some problems have polynomial-time approximation schemes (PTAS)

- **Parameterized Complexity:**

- Analyze complexity in terms of input size and additional parameter
- Example:  $O(2^k n)$  for vertex cover with parameter  $k$  = solution size
- Fixed-parameter tractable (FPT) if  $O(f(k) \cdot n^c)$

- **Heuristics and Metaheuristics:**

- No guarantees but work well in practice
- Genetic algorithms, simulated annealing, tabu search
- Often the only practical approach for large instances

## Algorithm Selection Guide

### Decision Framework

- **Understand the Problem Constraints:**

- Input size: Small ( $n < 50$ ), Medium ( $50 \leq n \leq 10^4$ ), Large ( $n > 10^4$ )
- Time constraints: Real-time, interactive, batch processing
- Space constraints: Memory-limited vs compute-limited
- Accuracy requirements: Exact vs approximate solutions

- **Analyze the Data Characteristics:**

- Sorted vs unsorted
- Random access vs sequential access
- Static vs dynamic (frequent updates)
- Distribution: Uniform, skewed, clustered

- **Consider the Operations Needed:**

- Search-heavy vs insert/delete-heavy
- Point queries vs range queries
- Need ordering or not
- Concurrent access requirements

## Quick Reference Table

Scenario	Recommended Approach	Time Complexity	Key Considerations
Small dataset	Quadratic sorts	$O(n^2)$	Simple, low constant factors
Large dataset sorting	Merge Sort / Quick Sort	$O(n \log n)$	Merge Sort stable, Quick Sort faster
Fast search, no ordering	Hash Table	$O(1)$ avg	No ordering, worst case $O(n)$
Ordered data, range queries	Balanced BST	$O(\log n)$	Maintains order, range operations
Optimization with optimal substructure	Dynamic Programming	Problem-dependent	Identify overlapping subproblems
Optimization with greedy choice	Greedy Algorithm	Usually $O(n \log n)$	Verify greedy choice property
Game playing	Minimax + Alpha-Beta	$O(b^{m/2})$	Evaluation function critical
Pathfinding with estimates	A* Search	$O(b^{ed})$	Need admissible heuristic
NP-complete problems	Approximation + Heuristics	Varies	Trade optimality for tractability
Geometric problems	Divide and Conquer	Often $O(n \log n)$	Exploit spatial properties

## Hybrid Approaches

- **Timsort:** Merge Sort + Insertion Sort - Python's built-in sort
- **Introsort:** Quick Sort + Heap Sort - C++ std::sort
- **B-trees:** BST + arrays - database indices, file systems
- **Bloom filters:** Hashing + probability - network routers, databases

## Final Principles: Developing Algorithmic Intuition

### The Mindset of an Algorithm Designer

- **Think in Trade-offs:** Every design decision has consequences. Understand what you're optimizing for and what you're sacrificing.
- **Embrace Asymptotic Thinking:** Focus on how algorithms scale. A "slow"  $O(n \log n)$  algorithm will eventually beat a "fast"  $O(n^2)$  algorithm.
- **Look for Patterns:** Most new problems resemble classic ones. Learn to recognize when to apply divide-and-conquer, dynamic programming, or greedy strategies.
- **Understand the Data:** The right data structure can make an intractable problem tractable. Choose based on the operations you need, not just what's familiar.
- **Prove Your Solutions:** Don't rely on intuition alone. Use mathematical reasoning to verify correctness and optimality.
- **Consider the Constants:** While asymptotic analysis is primary, constant factors matter in

practice. Profile and optimize hot paths.

- **Know the Limits:** Understand computational complexity theory. Recognize when you're facing an NP-hard problem and adjust your expectations accordingly.
- **Iterate and Refine:** Start with a simple solution, then optimize. Premature optimization is the root of much evil, but thoughtful optimization is the essence of good engineering.

## Continuous Learning

- **Study Classic Algorithms:** Understand why they work and their historical context
- **Analyze Real-world Systems:** Look at how algorithms are used in databases, operating systems, compilers
- **Practice Problem Solving:** Regular practice develops intuition and pattern recognition
- **Read Research Papers:** Stay current with new developments in algorithms
- **Implement and Experiment:** Theory informs practice, but practice deepens theoretical understanding